

# CSE 5523: HW3



# Outline

---

- You are to implement:
  - Linear logistic regression
  - Linear soft-margin SVM
  - Linear Naïve Bayes
  - Nonlinear Naïve Bayes

# Data

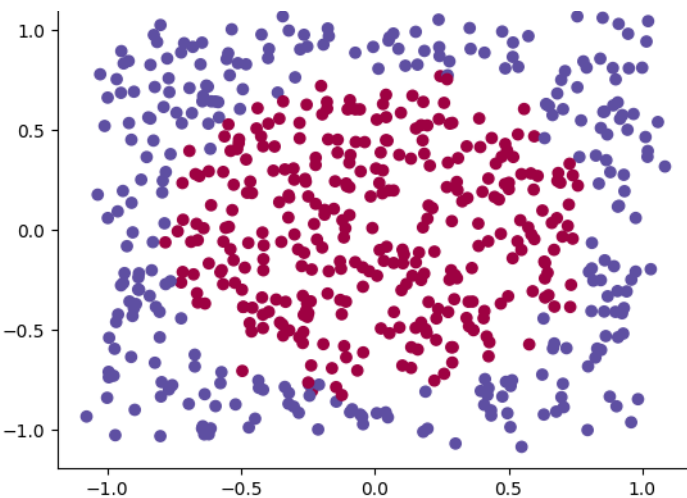
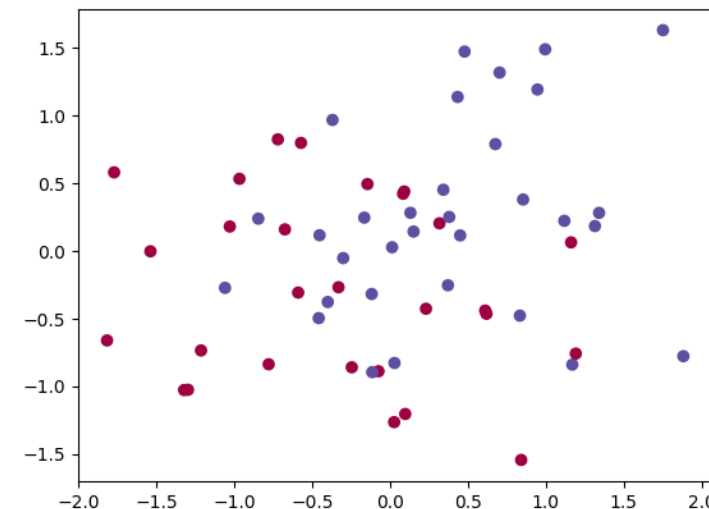
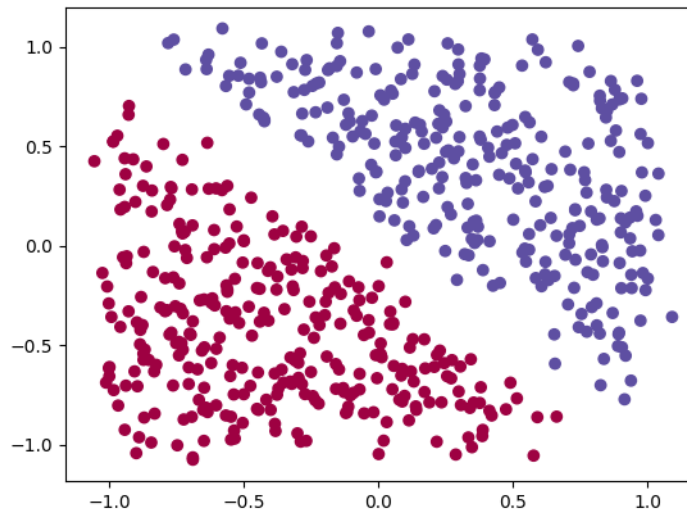
- Four data source

- 2D linear
- 2D noisy linear
- 2D quadratic (circle)
- MNIST (<5 vs. >=5)

- $X \in \mathbb{R}^{D \times N}$ :

- A column as an instance

- $Y \in \{+1, -1\}^{N \times 1}$



# Data

---

- The data  $\mathbf{X}$  are not appended with “1” yet.
- For feature transform for a 2D data instance  $\mathbf{x} \in \mathbb{R}^2$ , we do

$$\circ \boldsymbol{\phi}(\mathbf{x}) = \begin{bmatrix} x[1] \\ x[2] \\ x[1]^2 \\ x[2]^2 \\ x[1] \times x[2] \end{bmatrix}$$

- Again, you need to append “1” to the data  $\boldsymbol{\phi}(\mathbf{x})$  if you want to solve  $\tilde{\mathbf{w}}$  directly
- In the homework, we have done  $\boldsymbol{\phi}(\mathbf{x})$  for you!

# Accuracy

---

- Data =  $\{ (\mathbf{x}_i, y_i \in \{+1, -1\}) \}_{i=1}^N$
- Accuracy =  $\frac{1}{N} \sum_{i=1}^N \mathbf{1}[\hat{y}_i == y_i]$ , where  $\hat{y}_i$  is the prediction based on  $\mathbf{x}_i$

# Logistic regression



# Logistic regression

---

- **Training data:**  $D_{tr} = \left\{ \left( \mathbf{x}_i \in \mathbb{R}^D, y_i \in \{+1, -1\} \right) \right\}_{i=1}^N$
- **Model:**  $\text{sign}(\mathbf{w}\mathbf{x} + b) = \text{sign}(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}})$
- **Objective:**  $E(\tilde{\mathbf{w}}) = \frac{1}{N} \sum_{i=1}^N \log \left( 1 + e^{-y_i \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i} \right)$ 
  - Please add a  $\frac{1}{N}$  for normalization

# Gradient descent (GD) for logistic regression

---

- Initialize  $\tilde{\mathbf{w}}$
- For  $t = 1:T$ 
  - $\tilde{\mathbf{w}} \leftarrow \tilde{\mathbf{w}} - \eta \nabla_{\tilde{\mathbf{w}}} E(\tilde{\mathbf{w}})$  ;  $E(\tilde{\mathbf{w}})$ : the loss function you want to minimize!
  - No need to stop earlier
- Note:
  - If  $E(\tilde{\mathbf{w}}) = \frac{1}{N} \sum_n e_n(\tilde{\mathbf{w}})$ ; for example,  $e_n(\tilde{\mathbf{w}})$  is the loss on the  $i$ -th example
  - $\nabla_{\tilde{\mathbf{w}}} E(\tilde{\mathbf{w}}) = \frac{1}{N} \sum_n \nabla_{\tilde{\mathbf{w}}} e_n(\tilde{\mathbf{w}})$



# Gradient descent (GD) for logistic regression

---

- $e_n(\tilde{\mathbf{w}}) = \log \left( 1 + e^{-y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n} \right)$  for  $y_i \in \{+1, -1\}$

$$\begin{aligned}\nabla_{\tilde{\mathbf{w}}} e_n(\tilde{\mathbf{w}}) &= \frac{1}{1 + e^{-y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n}} \times \nabla_{\tilde{\mathbf{w}}} \left( 1 + e^{-y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n} \right) \\&= \frac{1}{1 + e^{-y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n}} \times (-y_n \tilde{\mathbf{x}}_n) \times e^{-y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n} \\&= \rho(y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n) \times (-y_n \tilde{\mathbf{x}}_n) \times e^{-y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n} \\&= \rho(y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n) \times (-y_n \tilde{\mathbf{x}}_n) \times \frac{1 - \rho(y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n)}{\rho(y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n)} \\&= -y_n \left( 1 - \rho(y_n \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n) \right) (\tilde{\mathbf{x}}_n)\end{aligned}$$

# Gradient descent (GD) for logistic regression

---

- $e_n(\tilde{\mathbf{w}}) = -[y_n \log p(+1|\mathbf{x}_n; \boldsymbol{\theta}) + (1 - y_n) \log(1 - p(+1|\mathbf{x}_n; \boldsymbol{\theta}))]$  for  $y_i \in \{+1, 0\}$

$$\begin{aligned}\nabla_{\tilde{\mathbf{w}}} e_n(\tilde{\mathbf{w}}) &= -[y_n \nabla_{\tilde{\mathbf{w}}} \log p(+1|\mathbf{x}_n; \boldsymbol{\theta}) + (1 - y_n) \nabla_{\tilde{\mathbf{w}}} \log(1 - p(+1|\mathbf{x}_n; \boldsymbol{\theta}))] \\ &= -\left[ y_n \nabla_{\tilde{\mathbf{w}}} \log \rho(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n) + (1 - y_n) \nabla_{\tilde{\mathbf{w}}} \log(1 - \rho(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_n)) \right] \\ &= -\left[ y_n \times (1 - \rho(\mathbf{w}^T \mathbf{x}_n)) \mathbf{x}_n - (1 - y_n) \times \rho(\mathbf{w}^T \mathbf{x}_n) \mathbf{x}_n \right] \\ &= -(y_n - \rho(\mathbf{w}^T \mathbf{x}_n)) \mathbf{x}_n\end{aligned}$$

# Soft-margin SVM algorithm



# Soft-margin SVM

---

- **Training data:**  $D_{tr} = \left\{ \left( \mathbf{x}_i \in \mathbb{R}^D, y_i \in \{+1, -1\} \right) \right\}_{i=1}^N$
- **Model:**  $\text{sign}(\mathbf{w}\mathbf{x} + b)$
- **Objective:**  $E(\mathbf{w}, b) = \frac{1}{N} \sum_n \max\{1 - y_n(\mathbf{w}^T \mathbf{x}_n + b), 0\} + \frac{1}{2} \lambda \mathbf{w}^T \mathbf{w}$ 
  - Please add a  $\frac{1}{N}$  for normalization
  - $\lambda$ : regularization coefficients (i.e., reg\_coeff in the code)

# Gradient descent (GD) for soft-margin SVM

---

- Initialize  $\tilde{\mathbf{w}}$
- For  $t = 1:T$ 
  - $\tilde{\mathbf{w}} \leftarrow \tilde{\mathbf{w}} - \eta \nabla_{\tilde{\mathbf{w}}} E(\tilde{\mathbf{w}})$  ;  $E(\tilde{\mathbf{w}})$ : the loss function you want to minimize!
  - No need to stop earlier
- Note:
  - If  $E(\mathbf{w}, b) = \frac{1}{N} \sum_n e_n(\mathbf{w}, b) + \frac{1}{2} \lambda \mathbf{w}^T \mathbf{w}$
  - $\nabla_{\mathbf{w}} E(\mathbf{w}, b) = \frac{1}{N} \sum_n \nabla_{\mathbf{w}} e_n(\mathbf{w}, b) + \lambda \mathbf{w}$
  - $\nabla_b E(\mathbf{w}, b) = \frac{1}{N} \sum_n \nabla_b e_n(\mathbf{w}, b)$

# Gradient descent (GD) for soft-margin SVM

---

To compute the gradients

- $\nabla_{\mathbf{w}} e_n(\mathbf{w}, b) = \begin{cases} -y_n \mathbf{x}_n, & \text{if } y_n(\mathbf{w}^T \mathbf{x}_n + b) < 1 \\ 0, & \text{otherwise} \end{cases}$
- $\nabla_b e_n(\mathbf{w}, b) = \begin{cases} -y_n, & \text{if } y_n(\mathbf{w}^T \mathbf{x}_n + b) < 1 \\ 0, & \text{otherwise} \end{cases}$

# Naïve Bayes



# Naïve Bayes

---

- **Training data:**  $D_{tr} = \{ (\mathbf{x}_i \in \mathbb{R}^D, y_i \in \{+1, -1\}) \}_{i=1}^N$
- **Goal:** construct  $p(Y = c|\mathbf{x})$  for  $\hat{y} = \max_{c \in \{+1, -1\}} p(Y = c|\mathbf{x})$
- **Bayes' rules:**  $p(Y = c|\mathbf{x}) \propto p(\mathbf{x}|Y = c)p(Y = c)$ 
  - $p(Y = c)$ : Bernoulli
  - $p(\mathbf{x}|Y = c) = p(x[1], x[2], x[3], \dots, x[D] | Y = c)$ 
$$= p(x[1] | Y = c)p(x[2] | Y = c) \dots p(x[D] | Y = c) = \prod_{d=1}^D p(x[d] | Y = c)$$
  - $p(x[d] | Y = c)$ : one-dimensional Gaussian (for all dimensions)



# Nonlinear Naïve Bayes

---

- $p(x[d] \mid Y = +1)$  and  $p(x[d] \mid Y = -1)$  have their own standard deviations  $\sigma_{d,+1}$  and  $\sigma_{d,-1}$
- See slides 11 or 12 for how to compute them
- You can represent  $[\sigma_{1,+1}, \sigma_{2,+1}, \dots, \sigma_{N,+1}]^T$  as a vector

# Linear Naïve Bayes

---

- $p(x[d] \mid Y = +1)$  and  $p(x[d] \mid Y = -1)$  share the same standard deviation  $\sigma_d$
- Built upon the previous slide, given  $\sigma_{d,+1}$ ,  $\sigma_{d,-1}$  and let  $N_{+1}$ ,  $N_{-1}$  be the number of training examples per class,  $\sigma_d^2 = \frac{N_{+1} \times \sigma_{d,+1}^2 + N_{-1} \times \sigma_{d,-1}^2}{N}$
- You can represent  $[\sigma_1, \sigma_2, \dots, \sigma_N]^T$  as a vector

# Prediction (please do “log” to prevent overflow)

---

$$\max_{c \in \{+1, -1\}} p(Y = c | \mathbf{x}) = \max_{c \in \{+1, -1\}} p(\mathbf{x} | Y = c) p(Y = c)$$

$$= \max_{c \in \{+1, -1\}} p(Y = c) \prod_{d=1}^D p(x[d] | Y = c)$$

$$= \max_{c \in \{+1, -1\}} \log p(Y = c) + \sum_{d=1}^D \log p(x[d] | Y = c)$$